

HOW TO AI UXR

Successfully integrating AI into the user experience research (UXR) workflow is the top objective on almost every research and ResearchOps professional’s mind.

“How to AI UXR” was produced to give you an overview of the key trends, help you identify the maturity level at which you’re implementing AI—crawl, walk, or run—and provide you with a list of implementations being used by research professionals across the globe, which you can apply too.

This resource was produced by Kate Towsey for *The ResearchOps Review* with the support of [Strella](#), the AI-powered customer research platform.

The Methodology

This map was produced using 562 data points gathered between January and May 2026. Data was gathered via desk research, a variety of one-on-one interactions, and during three working sessions involving fifty research and ResearchOps professionals proactively exploring the boundaries of AI-augmented research.

Human-produced content is a brand value for *The ResearchOps Review*, not because we’re Luddites, but because AI cannot (yet) produce original writing of the depth and rigour we’re committed to publishing. However, AI has been used to support the analysis and synthesis of this content—given that the production is about using AI to augment research, it would have been hypocritical not to use it. As many of the contributors in this production will know, AI enabled a scale and speed of analysis and synthesis that was previously impossible, but the process required systematic care to ensure that depth and rigour were preserved.

Understanding the Map

Most research professionals agree that the notion of a “standard research workflow” has always been inaccurate—studies vary in countless ways and rarely progress linearly—and AI is exploding the workflow as we know it. Even so, the contributors to this production agreed that a linear research workflow is still a useful, familiar framework on which to pin this novel content.

We’ve used the “crawl, walk, run” framework to provide you with a maturity matrix for implementing AI across the research workflow. As you view the map, you’ll notice that steps that were once discrete in the research workflow are now merging:

At the Crawl level, off-the-shelf AI tools are used to augment individual tasks for individual people, and the standard research workflow remains largely intact. AI is simply automating what used to be done by a human.

At the Walk level, the focus shifts from individual efficiency to team- or organisation-level systems. More than automating existing tasks, AI is customised to deliver previously impossible efficiencies. Also, parts of the research workflow are starting to merge, particularly in the analysis, synthesis, and insights packaging phases.

At the Run level, AI fundamentally alters how research operates. Research is delivered via human-in-the-loop (HITL) agentic systems, making it proactive rather than reactive. In several places, the sequential steps of a standard research workflow merge into “black box” steps, in which the reasoning that produced the insights is opaque and requires systematic evaluation (or “evals”).

The map doesn’t include a section dedicated specifically to ResearchOps, and for good reason. Because AI is a technology, the entire AI-augmented research workflow is being systematised—it’s all operations. That said, tasks that are typically considered operational are marked with a key.

The Explosion of Moore's Law

AI capabilities are growing at an accelerated rate, with agents’ task-solving abilities doubling roughly every seven months—far faster than the 24-month cycle of Moore’s Law. This map is accurate at the time of publishing, but just as AI is evolving, so are research technology vendors, the wider technology space, and the solutions of AI experimenters, like those who took part in this production. This is a snapshot in time, so explore the solutions with curiosity and care.

Contributors

Special thanks to the following AI innovators for their contributions:

Adam Valerio, Allison Robins, Angelica Eling, Anette Petersen, Anshuk Chhibber, Arev Pivazyan, Austen Lazarus, Aya Abdelgawad, Brian Greene, Brooke Sykes, Carina Cook, Caroline Cox-Orrrell, Casey Gollan, Christen Penny, Corina Kesler, Diana Sapanaro, Dr Asma Qureshi, Emily DiLeo, Farah Faisal, Filip Uzarevic, Graham Gardner, Hannah Mattil, Heidi Austin, Jordan Brinkman, Kalee Dankner, Kathy Shi, Kerttu Sobag, Katie Roehrick, Lydia Iana, Madeline Winer, Marshall Baker, Michel Vogel, Naki Ossom, Nathan Pena, Nicole Hack, Rachel Wigen-Toccalino, Rebecca Klee, Rita Cassilas, Shane Melton, Sohvi Silius, Stephanie Kingston, Stephanie Pratt, Tamia Sheldon, Tarah Srethwatanakul, Theresa Flood, Uyhun Ung

Let *The ResearchOps Review* know about your experiments. Mention [#HowtoAIUXR](#) and [@theresearchopsreview](#) on LinkedIn.

CRAWL

Crawl is about augmenting and automating individual tasks that would traditionally be done by humans. Off-the-shelf LLMs are being used to draft, summarise, cluster, and package across the research workflow, and to serve as a thinking partner. The work isn't systematised, and gains are individual and largely unseen at the organisational level, but a significant shift is already under way: steps that used to happen sequentially are being compressed into a single step or "conversation" with an AI.

- Insights Retrieval & Insights Creation Merge
- Democratisation
- ResearchOps
- The Emergent Make Phase
- AI Moderation & Synthetic Personas

Sponsored by

Strella

01

Prioritisation & Roadmapping

AI is helping teams to cluster requests, size issues, and surface knowledge gaps, but its tendency to affirm rather than interrogate means final prioritisation still calls for a researcher's judgement.

02

Method Selection & Scoping

AI is serving as a one-on-one sparring partner for method selection and study design. It's being used task-by-task, but without any underlying system or shared infrastructure to enable accuracy.

03

Existing Insights Retrieval & Reuse

LLMs, repositories, and wikis with built-in AI features are accelerating access to existing insights and enabling instant literature reviews. Researchers are also relying on deep-research tools and their citation features to verify quality.

04

Research Artefact Creation

AI is conquering the blank page, but outputs typically need significant editing before they're fit for purpose. Also, researchers are spending increasing time helping non-researchers fix work that looks polished but isn't.

05

Participant Recruitment

Recruitment is getting faster, but as AI-assisted responses become more common, spotting disingenuous or AI-generated data is becoming part of the recruitment workflow, too.

06

Data Collection

Data collection remains unchanged at this maturity level, but AI is automating real-time capture and analysis while extending what's possible in unmoderated (and now AI-moderated) research.

07

Data Preparation

AI is handling admin-heavy preparation tasks with speed and accuracy, and a significant shift is under way: data preparation is merging with analysis, synthesis, and packaging.

08

Analysis

Deep-research tools are transforming researchers' work by accelerating the data-to-insights cycle. Results are typically supplementary rather than complete and must be verified.

09

Synthesis

AI is being used as cognitive support to challenge and refine synthesis, provoke additional lines of inquiry, and compress data preparation to packaging into a single step once outputs have been evaluated.

10

Insights Packaging & Communication

AI is not only accelerating research content generation, it's also enabling researchers to build sharper reports and create multimedia formats that would previously have been impossible to deliver.

What Question Should We Answer?

Use AI to define which research questions should be answered and how to ensure teams tackle the most important ones.

Identify Knowledge Gaps

Use AI as a brainstorming tool to identify gaps and a study direction before formal conversations with stakeholders begin.

Organisational Topology

Use AI to identify key stakeholders, their roles and requirements for a research study, and create a chart that maps people to products to enable reliable filtering during the research intake phase.

Cut Down Long Lists

Use AI to cluster and prioritise numerous research requests based on prioritisation criteria, as a first step towards turning a long list into a sequenced, resourced research plan.

Brainstorming & Ideation Partner

Use an LLM as a sparring partner to refine research questions, brainstorm methodologies, challenge assumptions, and ideate on study structures.

Drafting & Rewording

Pull from existing templates and previous study data to assist in the creation of research plans and briefs, and rephrase business objectives into polished, machine-readable research questions.

Template Generation

Use AI to generate research templates to be used by researchers and non-researchers based on a study's requirements, budget, and deadline.

Automated Sizing of Issues

Use AI to organise and analyse massive volumes of customer data to accurately "size" an issue before committing to a study.

Source Interrogation

Ingest sources into a deep research tool to help understand the domain, interrogate documents, and identify which sources are worth reading in full.

Rapid Desk & Market Research

Use deep-research tools to quickly gather and synthesise external data and competitive reviews. Citation of sources is key to quality evaluation.

Multi-Source Summarisation

Group and summarise insights from disparate internal locations like SharePoint or Confluence, a research repository, customer feedback platforms, and more.

0-to-1 Drafting

Provide domain context (meeting recordings, past research, and internal documents) to an LLM to instantly draft research plans, screeners, and interview guides.

Contextual Material Creation

Use long-context tools like Google NotebookLM to generate workshop or focus group materials from background documents.

Synthetic Cognitive Testing

Stress-test interview activities with an LLM, instructed to take on a customer persona, to refine the plan. Also, use AI features in a survey tool to complete a survey with dummy data and see how the final visualisations will appear.

Job Mapping

Use jobs-to-be-done (JTBD) prompts to build structural maps of a user's workflow before starting discovery.

Outreach Communication Generation

Use an LLM to draft outreach emails, recruitment advertisements, and personalised greetings based on previous examples.

Clean Participant Lists

Use AI to extract names for personalisation, improving response rates: "Screener invitations with personalised greetings like 'Hello, Sarah' get a much better response rate."

Recruitment Strategy Brainstorming

Use AI as a sparring partner to develop a recruitment strategy, including where and how to find the target audience and how to communicate with them.

AI as Engagement Material

Use AI trends as a conversational hook to grow a participant panel. Everyone is interested in learning about AI.

Spot the AI

To support higher-quality recruiting, use AI to spot AI-like data-entry patterns and flag disingenuous responses. A human eye is still required to verify the results: some people just write well.

Real-Time Assistance

Use AI for note-taking and real-time summarisation during user interviews and usability tests.

Instant Session Roll-ups

Use AI to create summaries for individual sessions and "roll-ups" across multiple sessions. This can help identify patterns to explore or gaps in note-taking while research is ongoing.

Survey & Unmoderated Tool Features

Leverage built-in AI features in standard platforms for unmoderated tests.

Sensitive Topic Moderation

Use AI moderation for topics where participants might prefer more anonymity. Participants must always consent to how their data will be used.

Scorecarding in Minutes

Install a Figma plugin to streamline UX scorecarding, for instance, using the Measure UX Figma plugin.

PII Removal & Anonymisation

Redact personally identifiable information (PII) and instantly anonymise quotes from audio, text, or video recordings.

Transcript Formatting

Standardise formatting and rename files to align with organisational naming conventions, ensuring they're queryable in the future.

Automated Metadata Enrichment

Use AI to automatically categorise client attributes and enrich metadata tags, making insights significantly easier to retrieve from the repository later.

Data Clean-up

Quickly weed out low-quality data points, such as blank responses, untitled tests, or incomplete entries. Also, use AI to spot AI-like data-entry patterns and flag disingenuous responses.

Thematic Clustering & Coding

Identify surface-level themes in transcripts and create codebooks for large volumes of open-ended feedback. Compare and contrast the outputs with human-coded results to ensure accuracy, and fine-tune the LLM.

Annotating & Tagging

Create automated tagging of key moments, themes, tasks, and UI elements within transcripts to accelerate evidence retrieval.

Facilitate Analysis Workshops

Use AI to create facilitation plans for stakeholder insights workshops, enabling teams to work on analysis together.

Unmoderated Video Analysis

Use deep-research tools to watch session videos and describe what's happening. AI isn't yet able to read expressions or gestures, so this form of analysis should be considered supplementary.

Refining & Phrasing

Use AI to help reformat and adjust phrasings in findings or to find data that you remember but perhaps didn't code properly.

Counter-Bias Querying

Explicitly query the LLM for contradictions by asking it to find outliers, edge cases, and contradictory evidence to counterbalance and sense-check dominant themes.

ADHD Cognitive Support

Use AI-generated meeting summaries and action items as a supportive tool for ADHD researchers to manage the cognitive load of synthesis.

Did We Miss Anything?

Use AI to produce a gap analysis of the study's results and product briefs.

Rapid Presentation Building

Use tools like Gamma to build polished reports and decks; the AI will make suggestions about content density and visualisations.

Report Scaffolding & Phrasing

Use AI as a writing assistant to sharpen share-out language and help create templates with the target audience in mind.

Tailoring Insights for Impact

Manually remould findings for different audiences, such as turning research into executive reports, return-on-investment summaries, risk summaries, or engineering bug lists.

AI-Generated Multimedia Summaries

Create podcast-style audio briefings or video overviews that provide stakeholders with an engaging way to consume findings. The open question is whether audiences will tire of these novel content types.

Recycling & Roadmapping

Surface what's already known across different topics to prioritise where new research is genuinely needed and to spot knowledge gaps. "We use NotebookLM to load prior research and documentation for large mixed-methods work, map what's known, where evidence is inconsistent, and what questions remain."

What Do We Already Know?

Use deep-research tools to find key sources. Then vet each document's quality, print the best to PDF, and add them to a deep-research tool for interrogation and to support research artefact creation.

From Data to Insights

When research is low-risk, or AI outputs have been properly evaluated, AI is used to compress the final four stages of research into a single step: "We're using AI to structurally build and write compelling, actionable insights based on raw transcripts."

Transcript Clean-up & Summarisation

Use AI to clean up raw transcripts, remove filler words, and create initial summaries.

Rapid Cross-Case Analysis

Parse multiple interview transcripts into synthesis matrices to compare user responses to specific questions across multiple sessions.

Strategic Insight Weighting

Ask AI to provide specific theme strength, relevance, and reasoning for its ratings, moving beyond vague patterns to weighted evidence.

WALK

Walk is about building purpose-centred agents, skills, and Retrieval-Augmented Generation (RAG) systems that extend off-the-shelf AI. The focus shifts from individual efficiency to team-level systems, with evaluations, or "evals," emerging as a discipline in their own right. Synthetic personas and AI-moderated data collection create a need to understand the hidden processes behind "black-box insights" and guard against AI data loops. Finally, stages of the research workflow are starting to merge, particularly across analysis, synthesis, and insights packaging.

- Insights Retrieval & Insights Creation Merge
- Democratisation
- ResearchOps
- The Emergent Make Phase
- AI Moderation & Synthetic Personas

Sponsored by

Strella

01

Prioritisation & Roadmapping

Custom agents surface unvalidated assumptions, route research requests, and help non-researchers decide whether to self-serve or seek guidance before a researcher even enters the conversation.

02

Method Selection & Scoping

Context-trained agents help researchers and non-researchers move from a business question to a structured research plan, with the right method already built in.

03

Existing Insights Retrieval & Reuse

Custom RAG repositories and subject-matter-expert agents let partners query existing insights and "converse" with research without needing to switch tools or read full reports.

04

Research Artefact Creation

Purpose-built agents help non-researchers produce better artefacts and feel better prepared without needing guidance. Also, researchers can now build functional prototypes for usability testing.

05

Participant Recruitment

Custom agents build screeners, flag low-quality participants, manage panel logistics, and prepare participant dossiers, turning manual-heavy recruitment into an automated system.

06

Data Collection

Synthetic personas let teams pressure-test ideas before fielding studies with humans, and AI copilots handle note-taking as researchers revise the guide in real time.

07

Data Preparation

Agents flag low-effort responses and missed questions before analysis begins, parse transcripts into structured databases, absorb product flows, and translate raw data across languages.

08

Analysis

Analysis and synthesis are merging into a single loop: researchers now converse with the data through RAG-grounded chatbots, comparative journey maps surface diversity across user paths, and purpose-built pipelines theme the research and draft actionable insights. AI evals are increasingly crucial.

09

Synthesis

10

Insights Packaging & Communication

Reports are old news. Researchers are using rapid prototyping tools to show rather than tell. AI tools to create novel content that synthesises signals across multiple studies, and AI coding tools to deliver low-risk product changes.

■ First-Pass Triage

Build custom AI agents that enable stakeholders to do a first-pass triage without a researcher's help.

■ Automated Brief

Generation. Build a conversational agent that interviews designers about decisions and unknowns, then produces a structured research brief, significantly reducing vague research requests.

A Red-Flag Planner

Feed product documents into a prompt chain to surface unvalidated assumptions, flag user-sensitive decisions lacking evidence, and output a prioritised list of research opportunities.

■ Method Selection Assistance

Build AI assistants, such as a Custom GPT, to recommend specific methods, such as choosing between a survey, a usability test, or a diary study.

Informed Scoping

Use company- and project-specific contextual information, provided by agents that have access to internal data and a research repository, to make informed decisions about the optimal scope and research questions for a research project.

■ Instant Secondary Research Reports

Build agents to systematically collect, organise, and summarise research into instant secondary research reports.

Grounded RAG Repositories

Build a hierarchical research repository that includes master notebooks by customer segment, plus project-specific notebooks, to make user knowledge accurate, accessible, and grounded in sources.

■ Subject-Matter-Expert Agents

Use custom agents to serve as experts in specific product streams by pulling relevant data from team spaces. This enables stakeholders to "converse" rather than read full reports.

■ Stay Right Where You Are

Teams can query an agent directly from Slack to retrieve existing knowledge without switching platforms. This helps to keep the "conversation" with research insights going without interruptions.

■ Model Council Research

Use a model council (multiple LLMs) to do desk research on your behalf, then use another AI to synthesise the model council's reports to compare and contrast.

■ Strategic "Memos" from Signal Joins

Use custom prompts to gather internal data and findings across multiple studies into a single, synthesised "Research Strategic Memo."

■ A Synthetic Research Assistant

Build an AI agent trained on best practices to help researchers and non-researchers develop discussion guides, questionnaires, and recruitment screeners.

Context-Grounded Plan Generation

Use projects or custom agents loaded with deep domain context to ensure research plans respond to the context.

■ Automated Interviewer Coaching

Use specialised agents to review transcripts and provide feedback on interviewing skills based on the study's goals and UXR best practices.

■ Study Drafting

Use custom-built agents, like Custom GPTs and Gemini Gems, to help designers plan and build unmoderated studies.

■ Build It

Use dedicated tools like Claude Code or Codex to build functional prototypes for usability testing.

■ A Screener Question Expert

Build a skill that creates screener surveys using an extensive question library drawn from past studies. Use an LLM to build the question library.

■ Automated Fraud & Quality Detection

Use heuristics to flag low-quality participants. Use agents to watch for inconsistencies, overly generic or scripted responses, and to flag suspect applications or low-quality participants.

■ Automated Panel Management

Experiment with agents to handle administrative logistics such as reminders, rescheduling, and marking participant quality based on no-shows.

■ Participant Dossiers

Use agents to research and summarise participants' professional profiles and firmographics, and provide a summary prior to an interview.

■ Vet Potential Participants

Run AI-moderated interviews to figure out who would be most valuable for the team to talk to. Be aware not to waste potential participants' time, though, or risk brand integrity.

Analyst-Free Segmentation

Use AI-enabled analytics tools to more easily segment a target audience without the help of an analyst. This does require training the AI on data field names and descriptions.

■ "Vibe Check" Synthetic Personas

Build synthetic personas based on existing, well-researched personas. Use these to enable designers and PMs to "vibe check" their work before building the wrong thing. Also useful when the target audience is limited or hard to reach.

■ AI Heuristic Evaluators

Build simple apps using rapid prototyping tools to run quick heuristic evaluations of user interface screenshots before design critiques. "It can catch simple things we sometimes miss."

■ Hybrid Moderation Models

Design studies where an AI moderator hosts the interview, but participants have an explicit option to have a human join.

An Interview Copilot

Use AI to take notes, summarise research sessions in real time, allowing you to ask the AI questions and revise the interview on the fly.

■ Co-Creation Is Back

Run co-creation sessions with research participants in rapid AI prototyping tools. As an aside, these tools mean researchers are now doing more usability tests with interactive prototypes rather than mockups.

Behavioural Analysis Agents

Use specialised AI-powered assistants, like a HEX agent, to perform behavioural analysis on user logs and data to identify insights that traditionally required separate quantitative support.

Scale International Research

Use AI moderation tools to enable research in other languages. This used to be expensive and difficult, but AI moderation tools make it much more feasible.

Makeshift Database Parsing

Use AI to parse raw transcripts into structured databases or comparative journey maps during the preparatory phase.

Agents Thrive on Visual Context

For usability testing, provide a screen-by-screen description of product flows to support the quality of the analysis. AI performs better when it understands the product in detail.

Unlock Global Studies

Use AI formulas, say in Google Sheets, to instantly translate and interpret raw data.

Columns, Rows & Labels

For AI-assisted large-scale data analysis, say of surveys or NPS data, make data machine-ready by ensuring columns, rows, and labels are clear and well organised.

Converse with Data

Analysis is shifting from researchers analysing the data to using AI to engage in a back-and-forth "conversation" with data after conducting interviews, helping them to develop a deeper understanding of their users.

Ground Chatbots in Data

Create personalised chatbots using a RAG model, such as NotebookLM, to ground outputs in recordings and text documents prepared to assist with synthesis.

Analyse Massive Data Sets

Analyse data sets of open-ended feedback, such as NPS; use AI to create codebooks, then prompt it to code the data. To ensure accuracy, regularly evaluate outputs.

Confidence Scoring Across Signals

Use deep-research AI tools like Perplexity to combine multiple signals, such as qualitative research, product feedback, and sales notes, into a confidence score for an insight.

AI-Augmented Synthesis

Parse transcripts into a synthesis matrix for comparing users' answers to research questions. Use this to support AI-assisted, human-led analysis and synthesis.

■ Intake Recommendation Skills

Create skills that ask a series of questions about the project and provide a recommendation on how to do the research, whether a partner should self-serve or reach out to a UX researcher for guidance, and whether to do the research at all.

■ Synthetic Prioritisation & Scoping

Use agents trained on customer archetypes to support roadmapping and research scoping.

Predictive Gap Analysis

Provide product requirement briefs (PRDs) and transcripts to an agent to automatically identify inconsistencies in the evidence and unanswered questions. Use these to prompt further research, if necessary.

Machine-Readable Methodology

Convert established qualitative methodologies into machine-readable grounding documents to ensure AI's reasoning follows rigorous standards rather than pattern-matching.

Meeting to Research Plan to Artefacts in Seconds

Use skills to generate research plans from planning meetings, existing insights, recordings or transcripts, then generate discussion guides, a recruitment strategy, and screener surveys from existing question banks and templates.

■ A Synthetic Persona Library

Build AI personas based on actual customer data to explore the day-in-the-life, highlight gaps, compare personas, and improve research visibility beyond static documents.

■ Persona Chatbots & Simulators

Use agents trained on research interview data to build a qualitative preference-test simulator that can run tests and output a too-long-don't-read (TLDR) report.

■ A Synthetic Panel

Build a synthetic panel to simulate difficult-to-reach users or for low-risk research.

■ A Complete Agentic Pipeline

Build a purposefully designed pipeline of multiple agents to synthesise and theme research, then to write compelling, actionable insights.

RUN

Run changes the shape of a research practice entirely. The focus is production-grade research systems designed to anticipate organisational needs and push insights to stakeholders, rather than respond to a brief. Multi-agent and multi-step pipelines, automated workflows, and analytics integrations that monitor for signals across the business mean that the standard research workflow is condensing into far fewer steps. Because the reasoning behind AI outputs is often invisible, building eval processes to interrogate output quality is becoming a key operational requirement.

- Insights Retrieval & Insights Creation Merge
- Democratisation
- ResearchOps
- The Emergent Make Phase
- AI Moderation & Synthetic Personas

Sponsored by

Strella



01

Prioritisation & Roadmapping

Research roadmapping anticipates demand rather than reacts to it; agents track initiatives across the organisation, process large volumes of customer data, and surface gaps.

02

Method Selection & Scoping

Partners no longer need a researcher to scope a study; agents route requests to the right method, contact, or self-serve template, reserving researcher time for higher-priority initiatives.

03

Existing Insights Retrieval & Reuse

Insights retrieval now triggers instant content creation: agents draw on data warehouses, strategy documents, and internal knowledge to generate contextualised responses, and push relevant insights to stakeholders.

04

Research Artefact Creation

Artefact creation is now pipeline-driven: agents take a meeting transcript and produce a research plan, screener, and discussion guide, while multi-agent systems carry findings all the way through to a published output.

05

Participant Recruitment

Direct connections to data warehouses and in-product analytics enable pinpoint participant identification; agents then ease recruitment logistics, enabling ResearchOps to focus on system design.

06

Data Collection

The researcher's role is shifting from moderator to agentic system designers, and ResearchOps professionals can no longer work from the sidelines; to design these cognitive systems, they must have a clear understanding of the research craft.

07

Data Preparation

Pipelines now handle the full preparation workflow with minimal human intervention: removing PII, enriching metadata, classifying feedback, and staging files for analysis.

08

Analysis

Multi-agent pipelines blur the boundary between analysis, synthesis, and packaging; one agent extracts findings, another audits for gaps, and another writes the report and creates multimedia content that can be simultaneously posted to a repository, messaging platform, or wiki.

09

Synthesis

10

Insights Packaging & Communication

Insights retrieval and packaging are converging into a single activity; agents push personalised summaries to stakeholders based on their current work. The risk is that partners come to treat AI as an answer machine, reducing the demand for original research.

Data-Informed Roadmapping

Use AI to process and prioritise large amounts of customer data and feedback to inform the research roadmap.

Proactive Meeting Interventions

Build agents that "listen in" to stakeholder meetings and chime in with existing insights or suggest new research initiatives.

Methodological Advice

Build an agent trained on the standard operating procedure (SOP) for research and methodology to help non-researchers define the best approach to a study (assuming researchers aren't available to support).

Ethics & Privacy

Converse with an agent trained on research ethics, privacy, and risk to get feedback on study design.

AI-Ready Content Pipelines

Automate the conversion of static report documents into machine-readable formats, such as .md, for better machine ingestion.

Semantic & Vector-Based Search

Move beyond keyword matching to a search layer that retrieves insights by meaning and intent.

Managed Repository Agents

Build agentic systems that allow colleagues to "chat" with the research while reducing hallucinated findings because search queries are routed through custom agents.

Quality In. Quality Out.

Build an agent and skill to cross-reference data stored across multiple locations and to summarise insights. To support data integrity, the agent and the skill are instructed to use validated research insights rather than the raw data.

Ideas About IDEs

If research reports are stored in multiple locations, explore using AI Integrated Development Environments (IDEs) to make research documents accessible in a format that an agent or skill can digest. This may require automation to convert reports into machine-readable Markdown (.md) files.

Submit Your Research

Build browser extensions to streamline the submission of qualitative outputs to the research library, removing the "it's too manual" excuse.

Warehouse & BI Connectivity

Directly link AI to internal databases via MCP to use multiple data sources to quickly identify target participants who meet specific criteria.

Behavioural Targeting via Analytics

Couple in-product tools like Pendo, Hotjar, or Amplitude with LLMs for in-product targeting and pinpoint recruiting. This enables access for participants who are guaranteed to be users of a specific feature or workflow.

Red Tape Navigation

Build complex skills that internalise legal and go-to-market constraints, helping those doing research to identify high-signal participants and recruit within legal and privacy requirements.

Recruitment Platform Automation

Connect an agentic AI coding tool, such as Claude Code, to your recruitment platform to automate participant recruitment. Input: describe the type of users you need. Output: the study is set up with the relevant screeners.

AI Panel Manager

Build an AI system that manages an internal enterprise panel. It can identify participants, track participation and cool-down periods, and provide partner-facing cohort summaries without exposing personally identifiable information (PII).

Weekly Panel Health Reports

Build an MCP-enabled agent to provide a weekly report on population gaps and health analysis to support research participant panel management.

Qualitative Research at Quantitative Scales

Run ~100 asynchronous AI-moderated sessions per week using tools like Strella to create a continuous "insight stream."

Agentic Feedback Loops

Set up n8n automations (a low-code workflow automation tool) connected to Slackbots and Gemini to ask customers dynamic follow-up questions and automatically classify feedback.

Automated Feedback Ingestion & Reporting

Set up pipelines to automatically tag and summarise customer calls, feedback, and complaints and push them directly to Slack channels, enabling team members to review feedback at a glance. Automatically deliver a weekly summary of top issues.

Automated Innovation Labs

Support massive research events (e.g., 150+ sessions) by using AI to enhance moderator notes and check drafts against original recordings for accuracy.

Dynamic Moderated Probing

Implement purposeful AI moderation that automatically requests follow-ups and probes for behavioural depth during churn or post-purchase interviews.

Rapid Timelines & Anonymity

Use AI moderation tools to support rapid timelines or when participants might prefer greater anonymity, e.g., in research on fraud. As with human-led research, participants must still consent to what data will be captured and how it will be used.

A Metadata Machine

Use AI to analyse interview transcripts and automatically enrich metadata and categorise client attributes to make insights easier to find later.

Quantitative Assistance

Generate Python code for statistical analysis (e.g., clustering or regression). This requires the researcher to have a strong foundation in statistics to verify the output.

Scaled-up Qualitative Research

Use tools like Google Sheets plus Gemini or Sheets AI formulas to filter and process hundreds of open-ended responses into mental models.

Proactive "Push" Delivery

Agents proactively deliver personalised insights to stakeholders based on their current activity and interests, e.g., internal documents they look at the most, specific product metrics, or projects they're working on.

Efficacy & Alignment Intelligence

Build idea-scoring skills that assign quantitative and qualitative scores to the efficacy of an idea against corporate and team plans. The agent also coaches partners on where there's an overlap with other product areas and an opportunity to align.

Automatic Prototype Creation

Experiment with training AI coding tools to automatically create or iterate on a prototype in response to a research report.

Natural Language Roadmapping to Recruiting

Connect an agent to your data warehouse that identifies participants based on natural language requests. This is one of many skills programmed into the agent; recruitment support is always offered as the next step after supporting the user through prior steps in the process. Agree on guardrails and limitations with customer support managers (CSMs) on customers who can and can't be recruited.

Analytics & LLM-Coupled Scoping

Combine customer analytics data from tools like Pendo with LLM-based research analysis to identify churn and problem areas, and automatically scope and plan research to address those issues. Then use Pendo Guides or Lightbox (customer communication tactics) to identify potential research participants experiencing those issues.

Tools That Network & Track

Build planning and customer recruitment skills that include tracking, so that, as stakeholders use them, the research team can see initiatives across the team to build research roadmaps around. "We've built entire tools that our product design team uses to plan and map work across the team. It includes designer profiles, so you can more easily network with the right designers for feedback."

Integrated Operational Support

Build agents to support a product stream that not only use the repository as a research subject-matter expert but also pull related data from product team spaces to generate actions, summaries, Jira tickets, etc.

A Research Repository & Research Assistant

The repository houses hundreds of research studies: UXR studies, competitive research, social listening reports, behavioural analytics, and NPS and brand tracking data. Partners are encouraged to ask the repository questions, identify gaps before running new research, and then work with the AI assistant to scope the study, create a research brief, and produce research artefacts.

Transcript-to-Script Pipelines

Leverage AI and specialised prompt libraries to listen to planning meetings, then to automatically generate 80-percent-complete research plans and screeners based on the transcript; from transcripts, to prompt library, to human-in-the-loop (HITL), to a test script.

Surface Additional Context

Build an agent that runs a keyword search in the background to surface related secondary research when someone asks a related question or mentions starting a new research study. This helps to remove duplicate work, surface existing research to add more context, and build on existing insights.

A "Which Method? Routing Agent

Build complex systems that route partners who want to do research down specific paths. The agent should advise them which methodology to use, when a methodology is too complex for self-service, and whom to contact for additional support.

Template-Compliant Automation

Create agentic skills that follow organisational templates and automatically store results. "It takes a planning meeting Zoom recording and turns it into a draft research plan that follows our planning template, and saves it to Confluence via MCP."

Risk Assessment

Build an agentic framework that uses behavioural analytics to ascertain when traditional research steps are crucial and when they can be skipped. For high-risk research, pull out all the stops. But for low-risk, low-priority research, research can be handled via AI moderation, self-service, or by going straight to production.

Nothing in the library?

If a search doesn't surface content in the library, train the agent to reroute the searcher to request original research from the research team. Use AI to sort requests by strategic alignment to inform the research roadmap.

Continuous "Insight Streams"

Use a custom agent to regularly scrape and synthesise internal data (NPS, support tickets, sales notes, customer call transcripts, and in-app queries) and external data (Reddit threads, App Store reviews) and publish weekly digests, or real-time summaries directly into Slack, on broader customer sentiment beyond research studies. This requires consistent tagging and data governance.

Instant Weekly PESTLE Briefings

Build a weekly trigger in a desktop AI agent, such as Claude Cowork, to run concurrent PESTLE (political, economic, social, technological, legal, and environmental signals) queries, log the results in a database, and deliver weekly briefings to partners. This provides partners with a regular update on external market forces.

Hands-off Research

Conduct AI-moderated interviews on topics such as customer churn, post-purchase interviews, to A/B test qualitative insights, or to explore post-launch feedback. AI moderation tools can independently create the study guide and recruitment screener and launch the study. Researchers use AI moderation as part of an end-to-end AI workflow when research is low-risk.

An End-to-End Reporting Agent

"I created an agentic system with Claude Code: one agent extracts findings from interviews, another generates insights, another checks interviews for missing evidence, and another checks for alternative interpretations. Another agent writes a report and learns my tone of voice; another can create presentations."

An Automated Analysis

Synthesis, and Share-Out Pipeline. Use an AI code-writing tool to build a Python script that monitors a folder; when transcript files are added to the folder, an AI agent removes PII from and cleans up the transcripts, adds metadata tags, renames the files based on criteria, logs updates, writes an interview summary, and then moves the files to a "ready to upload" folder.

Agentic Report Writing Systems

Build multi-agent systems (e.g., via Claude Code) that extract findings, generate insights, audit for missing evidence, and write reports that learn a researcher's specific voice and a stakeholder's specific needs.

A Multi-Agent Qualitative Analysis & Synthesis Pipeline

Automate analysis and synthesis by building an agentic system featuring agents trained on empathy-based coding methodology and each deployed in a functional role: reasoning, execution, and auditing. For instance, one agent extracts findings while another checks for missing evidence or alternative interpretations.

GLOSSARY

AI Evaluation or “Eval”

The systematic process of measuring an AI model's or agent's performance, accuracy, safety, and reliability against predefined criteria; without evaluation, subjective judgment becomes the acceptance test, and quality degrades quietly over time.

AI Integrated Development Environment (IDE)

An integrated development environment (IDE) with AI capabilities that supports code-aware chat, multi-file edits, and debugging; these tools make rapid prototyping materially easier for non-engineers, which changes how quickly research teams can ship artefacts.

AI-Moderated Research

Data collection in which an AI system conducts part or all of a research session (for example, asking questions, probing, and following up), sometimes alongside a human moderator. AI moderation can increase speed and volume, but it raises quality, consent, and “black-box” visibility risks if poorly designed.

Agent

An autonomous, goal-oriented AI system that can plan steps, make decisions, and take actions to complete a task. Agents can compress multi-step workflows into “black-box” output unless you design for traceability, review, and governance.

Agentic Workflow

A workflow in which one or more agents coordinate tasks end-to-end (for example, collecting inputs, cleaning data, synthesising findings, drafting artefacts, and routing work for approval); it changes the unit of work from “task assistance” to “workflow ownership,” which raises risk, QA, and accountability questions.

Bias Amplification

When AI outputs over-index on dominant, surface-level themes, flattening nuance and making minority signals harder to see, creating false confidence in the “main story,” especially when stakeholders want quick answers.

Confidence Scoring

A structured estimate of how strongly an insight is supported by available evidence, often by combining multiple signals. Confidence scores can create false precision unless they are tied to explicit criteria and traceable sources.

Context Window

The amount of text (measured in tokens) a model can consider at once; long-context models can process very large inputs, from ~128k to 1M+ tokens, but they don't guarantee the model will attend to the right evidence.

Custom GPT

A customised, pre-configured version of a ChatGPT model that includes specific instructions, tools, or knowledge sources; it can standardise work across a team, but it can also standardise mistakes if not reviewed and maintained.

Data Loop

A situation in which AI-generated outputs feed back into the system as inputs, shaping future outputs. Data loops can quietly reduce contact with real customer evidence unless you design for grounding and periodic reality checks.

Deep-Research Tools

AI-powered agents that autonomously browse, analyse, and synthesise information from many sources into a report; they can be fast and useful, but they also concentrate citation risk and require careful verification.

Embeddings

Numerical representations of text that capture semantic meaning and enable similarity search; embeddings influence what a retrieval system “finds,” and therefore what it can credibly say.

Fine-Tuning

Training a model further on a targeted dataset to improve performance on a specific type of task. Fine-tuning can improve consistency, but it can also bake in outdated assumptions if the underlying data is stale.

Grounding

Constraining an AI system's output to verified sources (for example, retrieved documents, transcripts, or approved policy); grounding is a practical antidote to confident but untraceable output.

Hallucination

When a model produces information that is incorrect, fabricated, or not supported by the available evidence, often with high confidence. In research, hallucinations aren't always “minor errors;” they can lead to strategy and product decisions or public claims.

Human-in-the-Loop (HITL)

A system-design approach in which a person must review, approve, or correct AI outputs at particular points before it's finalised or acted upon, putting accountability back into the system, especially for high-risk decisions.

Long-Context Tools

Long-context AI models and frameworks are designed to process very large volumes of information in one pass, often 128k to 1M+ tokens; they reduce summarisation overhead, but they still require good input hygiene, including redaction, document selection, and metadata.

Model Council

An AI model council is a multi-model research architecture. It runs a single query through several AI models simultaneously to provide a unified, cross-verified response.

Multimodal

Models that can work with more than text, such as images, audio, and video, have multimodal capabilities that change what “data” can be analysed and raise new consent and privacy issues.

Prompt Chaining

A technique in which a complex task is broken into smaller subtasks, completed across multiple prompts, with intermediate outputs feeding the next step; it improves structure and reliability, but it also multiplies points of failure unless you add checks.

Prompt Engineering

The craft of writing instructions that reliably produce useful output, including specifying constraints, inputs, format, and evaluation criteria. It's a lightweight workflow design, not a matter of “magic wording.”

Prompt Library

A curated set of tested prompts and prompt chains used repeatedly for common tasks. A prompt library reduces rework and drift, but it needs version control and regular review as tools and contexts change.

Prompt Template

A reusable prompt structure with fixed sections, such as inputs, constraints, output format, and checks; templates make output more consistent and easier to evaluate.

Retrieval-Augmented Generation (RAG)

An AI framework that improves the accuracy and reliability of large language models (LLMs) by giving them access to external, authoritative data sources before they generate a response. RAG systems can make answers more trustworthy, but only if retrieval is good and the system preserves provenance via citations.

Skills

Reusable, structured procedures that an agent or user can apply to perform a specific task; they standardise quality across a team and make it easier to audit how work is done.

Synthetic Data

Data generated by AI rather than collected from real people; it can be useful for prototyping and stress-testing, but it can also create circularity and false confidence if treated as empirical evidence.

Temperature

A model setting that influences randomness in generation; for research artefacts, reliability usually beats creativity, and temperature is a hidden quality lever.

Tokens

The fundamental, atomic unit of data that an LLM processes. Instead of reading sentences as a continuous stream like humans do, AI models break text down into these discrete “chunks” to analyze patterns and generate responses. Token budgets and management are an emerging and important operational theme.

Vector Database

A database optimised for storing and querying embeddings for semantic retrieval; if your system retrieves evidence, the vector database is effectively its memory.

Vector-Based Search

A technique that finds information based on semantic meaning and conceptual similarity rather than exact keyword matches, typically using embeddings. Vector-based search improves discovery, but it can also return plausibly related content that is not actually relevant unless tuned and evaluated.

Verification

The practice of checking AI outputs against original sources, including quotes, numbers, claims, and attribution; it is the operational antidote to hallucination and the strategic antidote to reputational risk.